

The Four Recipes

Take-home one-pager. Transferable to any brand, any dataset, any AI tool.

Recipe 1. Start from the business question, not the data.

Bad first move: "Find themes in these reviews."

Better first move: "What is Starbucks' VP of Mobile Experience worried about this quarter? Which themes would inform that?"

Every dashboard, every chart, every theme should serve a named decision-maker and a named question. If you can't say which exec your slide helps, cut the slide.

Try it: write three exec questions before you write a single prompt. The questions are your north star for the next 50 minutes.

Recipe 2. Bridge qual to quant with a codebook, not vibes.

The naive AI workflow is "show the AI the reviews and ask for patterns." It will give you something. It will not be defensible.

The real workflow:

1. Draft a codebook (5-10 named themes with definitions and example phrases)
2. Apply the codebook to every review programmatically (one prompt per review, structured JSON out)
3. Now you have a structured table: review_id, stars, city, themes[]
4. Now you can chart, filter, and pivot like quant data

A codebook is a contract between you and the AI. Without it, the AI is guessing what you care about. With it, you can audit, refine, and compare runs.

Try it: extend the starter codebook with one theme you noticed but it missed. Re-run coding on 50 reviews. Watch the chart change.

Recipe 3. Cite or it didn't happen.

For every theme assignment, the AI must return the verbatim quote span that supports it. No paraphrasing. No summaries. If the quote doesn't literally appear in the review text, throw it out.

Why: this is what makes your insight portfolio-grade. When the VP asks "really? show me," you point at the row. The quote is in the data. The data is from Yelp. There is no hallucination layer between your claim and the evidence.

This is the same principle Kai's [Patterns Lite](#) bakes in with a dedicated verifier subagent. We're doing the simpler version inline: filter out any quote that fails the substring check.

Try it: open three reviews tagged with the same theme. Read the supporting quotes. If any feels stretched, your codebook definition is too loose.

Recipe 4. The dashboard is a question machine.

- Filters are not decorations. Each filter is a follow-up question you anticipated the exec asking.
- City filter = "is this a [Tampa / Philly / Nashville] problem or a national one?"
 - Star rating filter = "is this concentrated in 1-2 star reviews, or are 4-star fans complaining too?"
 - Date filter = "is this getting worse or better?"
 - Theme filter = "which themes drive 1-star reviews most?"

If a filter doesn't map to a real question someone would ask, delete it.

The dashboard is also a question machine in a second sense: open it in Cursor agent mode and ask plain-English questions of the CSV ("what percent of mobile_app_order reviews are 1-2 stars?"). Cursor reads the file, runs Python, answers with chart or table. That's your live interview surface.

Try it: in the live Q&A, when the VP asks something your dashboard doesn't directly answer, switch to Cursor and have it run the answer. Show your work.

Putting it together



Each recipe is doing real work. Skip any one and the chain breaks: vibes-based themes (no codebook), uncited claims (no verifier), pretty charts without filters (no question machine), or filters and themes that don't connect to any business question.

What to take to your next portfolio piece

1. Pick a brand. (Any brand with public reviews works.)
2. Write three exec questions before touching the data.
3. Draft a starter codebook from the questions.
4. Code 500–2,000 reviews with structured–JSON output and quote citations.
5. Build a small dashboard with filters that map to the questions.
6. Write a 1–page insight memo where every claim is grounded in a quote.
7. Stress–test it: have a friend play exec and grill you for 5 minutes.

That's a portfolio piece you can defend.